

---

# Towards Interpretable Natural Language Understanding with Explanations as Latent Variables

---

Anonymous Author(s)

Affiliation

Address

email

## Abstract

1 Recently generating natural language explanations has shown very promising re-  
2 sults in not only offering interpretable explanations but also providing additional  
3 information and supervision for prediction. However, existing approaches usually  
4 require a large set of human annotated explanations for training while collecting  
5 a large set of explanations is not only time consuming but also expensive. In  
6 this paper, we develop a general framework for interpretable natural language  
7 understanding that requires only a small set of human annotated explanations for  
8 training. Our framework treats natural language explanations as latent variables  
9 that model the underlying reasoning process of a neural model. We develop a  
10 variational EM framework for optimization where an explanation generation mod-  
11 ule and an explanation-augmented prediction module are alternatively optimized  
12 and mutually enhance each other. Moreover, we further propose an explanation-  
13 based self-training method under this framework for semi-supervised learning.  
14 It alternates between assigning pseudo-labels to unlabeled data and generating  
15 new explanations to iteratively improve each other. Experiments on two natural  
16 language understanding tasks demonstrate that our framework can not only make  
17 effective predictions in both supervised and semi-supervised settings, but is also  
18 able to generate good natural language explanations.

## 19 1 Introduction

20 Building interpretable systems for natural language understanding is critical in various domains such  
21 as healthcare and finance. One promising direction is generating natural language explanations for  
22 prediction [1–4], which has been shown very promising recently as they are not only able to offer  
23 interpretable explanations for back-box prediction systems but also provide additional information  
24 and supervision for prediction [5–7]. For example, given a sentence “*The only thing more wonderful*  
25 *than the food is the service.*”, a human annotator may write an explanation like “*Positive, because the*  
26 *word ‘wonderful’ occurs within three words before the term food*”, which is much more informative  
27 than the label “*positive*” as it explains how the decision was made. Moreover, the explanation can  
28 serve as implicit logic rules that can be generalized to other instances like “*The food is wonderful, I*  
29 *really enjoyed it.*”

30 There are some recent work [3, 4] that study generating natural language explanations for predictions  
31 and/or leverage generated explanations as additional features for prediction. For example, Camburu  
32 et al. [3] trained a language model to general natural language explanations for the task of natural  
33 language inference by training on a corpus with annotated human explanations. Rajani et al. [4]  
34 proposed a two-stage framework for common sense reasoning which first trained a natural language  
35 explanation model and then further trained a prediction model with the generated explanation as  
36 additional information. These approaches achieve promising performance in terms of both prediction

and explainability. However, a large number of labeled examples with human explanations are required, which is expensive and sometimes impossible to obtain. Therefore, we are looking for an approach that makes effective prediction, offers good explainability, but requires a limited number of human explanations for training.

In this paper, we propose such an approach. We start from the intuition that the explanation-augmented prediction model is able to provide informative feedback for generating meaningful natural language explanations. Therefore, different from existing work which trains the explanation generation model and the explanation-augmented prediction model in separate stages, we propose to **jointly** train the two models. Specifically, taking the task of text classification as an example, we propose a principled probabilistic framework for text classification, where natural language Explanations are treated as **Latent Variables (ELV)**. Variational EM [8] is used for the optimization, and only a set of human explanations are required for guiding the explanation generation process. In the E-step, the explanation generation model is trained to approximate the ground truth explanations (for instances with annotated explanations) or guided by the explanation-augmentation module through posterior inference (for instances without annotated explanations); in the M-step, the explanation-augmented prediction model is trained with high-quality explanations sampled from the explanation generation model. The two modules mutually enhance each other. As human explanations can serve as implicit logic rules, they can be used for labeling unlabeled data. Therefore, we further extend our ELV framework to an **Explanation-based Self-Training (ELV-EST)** model for leveraging a large number of unlabeled data in the semi-supervised setting.

To summarize, in this paper we make the following contributions:

- We propose a principled probabilistic framework called ELV for text classification, in which natural language explanation is treated as a latent variable. It jointly trains an explanation generator and an explanation-augmented prediction model. Only a few annotated natural language explanations are required for guiding the natural language generation process.
- We further extend ELV for semi-supervised learning (the ELV-EST model), which leverages natural language explanation as implicit logic rules to label unlabeled data.
- We conduct extensive experiments on two tasks: relation extraction and sentiment analysis. Experimental results prove the effectiveness of our proposed approach in terms of both prediction and explainability in both supervised and semi-supervised setting.

## 2 Related Work

Natural language (NL) explanations have been proved very useful for both model explanations and prediction in a variety of tasks recently [9, 10, 3, 4, 4]. Some early work [11, 9, 10] exploited NL explanation as additional features for prediction. For example, Srivastava et al. [9] converted NL explanations into classifier features to train text classification models. Fidler et al. [12] used natural language explanations to assist in supervising an image captioning model. Very recently, Murty et al. [10] proposed ExpBERT to directly incorporate NL explanations with BERT. However, most of these work require the explanations to be available in both training and testing instances, which is not realistic as annotating the explanation of a huge amount of instances is very time consuming and expensive. Moreover, the prediction becomes much easier once the explanations are given in the testing data.

There is some recent work that studied training a natural language explanation model and then used the generated explanations for prediction. For example, Camburu et al. [3] and Rajani et al. [4] proposed to separately train a model to generate NL explanations and a classification model that takes the generated explanations as additional input. Their approaches have shown very promising for improving the interpretability of classification models and increasing the prediction performance with explanations as additional features. However, their approaches require a large number of human annotated NL explanations to train the explanation generation model. Moreover, these approaches fail to model the interaction between generating NL explanations and exploiting NL explanations for prediction. As a result, there is no guarantee that the generated explanations reflect the decision-making process of the prediction model or beneficial to the prediction model. As reported by Camburu et al. [3] that interpretability comes at the cost of loss in performance. In this paper, we propose a principled probabilistic framework with explanations as latent variables to minimize the number of

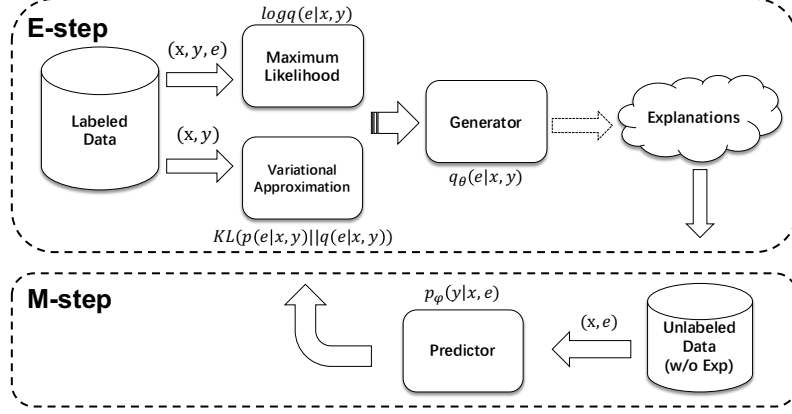


Figure 1: Overview of ELV. During E-step, we train our generator  $p(e|x, y)$  to generate explanations given labeled data. For labeled data with annotated explanations (i.e.  $\mathcal{D}_E$ ), we maximize the likelihood of the ground truth explanations. For labeled data without explanations (i.e.  $\mathcal{D}_L$ ), we minimize the KL divergence between the variational distribution  $q_\theta(e|x, y)$  and the ground truth posterior  $p(e|x, y)$ , which is calculated with the help of the prediction model. During M-step, we use the explanation generated in E-step to train the predictor  $p(y|x, e)$  with MLE.

90 training instances with explanations by jointly training the natural language explanation module and  
 91 the explanation-augmented prediction module.

92 Another relevant direction is treating natural language explanations as additional supervisions for  
 93 semi-supervised learning instead of as additional features [13, 7]. For example, Hancock et al. [13]  
 94 utilized a semantic parser to parse the NL explanations into logical forms (i.e., “labeling function”).  
 95 The labeling functions are then employed to match the unlabeled examples either hardly [13] or  
 96 softly [7] to generate pseudo-labeled datasets used for training models. However, these approaches  
 97 require the explanations to be annotated in a form that can be accurately parsed by a semantic  
 98 parser to form labeling functions, which may not be possible for many NLP applications. In our  
 99 semi-supervised framework, semantic parsing is not required, and natural language explanations  
 100 are interpreted with distributed representation obtained by pre-trained language models for labeling  
 101 unlabeled data.

## 102 3 Methodology

### 103 3.1 Problem Definition

104 Given an input sentence  $x$ , we aim to predict its label  $y$  and generate a natural language (NL)  
 105 explanation  $e$  describing why  $x$  is classified as  $y$ . Specifically, given a few amount of training  
 106 example with NL explanation annotation  $\mathcal{D}_E = \{(x_1, y_1, e_1), \dots, (x_n, y_n, e_n)\}$  and a relatively  
 107 large set of labeled examples  $\mathcal{D}_L = \{(x_{n+1}, y_{n+1}), \dots, (x_m, y_m)\}$ , our goal is to learn: 1) An  
 108 explanation generation model  $E_\theta$  that parametrize  $q(e|x, y)$ , which takes a labeled example  $(x, y)$   
 109 as input and generates a corresponding natural language explanation  $e$ , and 2) an explanation-  
 110 augmented prediction model  $M_\phi$ , which parametrize  $p(y|x, e)$  and takes an unlabeled example  $x$  and  
 111 NL explanations (as implicit rules)  $E$  to assign a label  $y$  to  $x$ .

### 112 3.2 Natural Language Explanation as Latent Variables

113 Given labeled data  $(x, y)$ , we treat the nature language explanation  $e$  as a latent variable. For training,  
 114 we aim to optimize the evidence lower bound (ELBO) of  $\log p(y|x)$ , which can be formulated as:

$$\log p(y|x) = \log \int p(e, y|x) de = \log \int q_\theta(e|x, y) \frac{p(e, y|x)}{q_\theta(e|x, y)} de \quad (1)$$

$$\geq \mathcal{L}(\theta, \phi) = E_{q_\theta(e|x, y)} \log \frac{p(e, y|x)}{q_\theta(e|x, y)} = E_{q_\theta(e|x, y)} \log \frac{p_\phi(y|e, x)p(e|x)}{q_\theta(e|x, y)}, \quad (2)$$

where  $q_\theta(e|x, y)$  is the variational distribution of posterior distribution  $p(e|x, y)$ ,  $p(e|x)$  is the prior distribution of explanation  $e$  for instance  $x$ , and  $p_\phi(y|e, x)$  is the explanation-augmented prediction model.

Due to the large search space of natural language explanation  $e$ , instead of using reparametrization trick used in variational autoencoder [14], here we use the variational-EM algorithm for optimization. The overview of ELV is illustrated in Figure 1. Note that in our training data, the explanations of a few number of labeled examples (i.e.,  $\mathcal{D}_E$ ) are given. Therefore, we first initialize the explanation generation model  $E_\theta = q_\theta(e|x, y)$  and the prediction model  $M_\phi = p_\phi(y|e, x)$  by training on  $\mathcal{D}_E$  with maximum likelihood estimation (MLE). Then, we update the aforementioned models by maximizing the log-likelihood  $\log p(y|x)$  using  $\mathcal{D}_L \cup \mathcal{D}_E$  with variational-EM. In the variational E-step, we train the explanation generator to minimize the KL divergence between  $q_\theta(e|x, y)$  and  $p(e|x, y)$ , which will be detailed in Section 3.3. In the M-step, we fix  $\theta$  and  $p(e|x)$  and update the parameters  $\phi$  of the prediction model to maximize the log-likelihood  $\log p(y|x)$ .

### 3.3 E-step: Explanation Generation model

As the core component of our variational EM framework, the explanation generation model is expected to generate “soft” logic rules in the form of natural language explanation. However, training a seq2seq model to generate explanations of high quality from scratch is very challenging. Motivated by the recent finding that pretrained language models encode various types of factual knowledge and commonsense knowledge in their parameters [15–17], we employ UniLM [18]—a unified pre-trained language generation model that achieved the state-of-the-art performance on many text generation tasks—as the explanation generation model  $E_\theta$  in our framework. Specifically, the explanation generation model takes as input the concatenation of input sentence  $x$  and the text description of its corresponding label  $y$  to generate an explanation, which explains the label decision in natural language and can be treated as an implicit logic rule that can generalize to other examples.

Note that in the training data, there is a small set of labeled examples with explanations and the others without explanations. Therefore, in the variational E-step, for labeled data  $(x, y)$  without explanations (i.e.  $\mathcal{D}_L$ ), we are trying to use the variational distribution  $q_\theta(e|x, y)$  to approximate the ground truth posterior  $p(e|x, y)$ , which can be calculated as

$$p(e|x, y) \sim p_\phi(y|x, e)p(e|x) \quad (3)$$

where  $p_\phi(y|x, e)$  is parameterized by the prediction model and provides feedback for generating meaningful natural language explanation. We will introduce the detailed parametrization of  $p(e|x)$  and  $p_\phi(y|x, e)$  in the M-step.

For labeled data with explanations (i.e.  $\mathcal{D}_E$ ), we just need to maximize the likelihood of the ground truth explanations. Therefore, the overall objective function of E-step can be summarized as:

$$O = \sum_{(x, y) \in \mathcal{D}_E} \log q(e|x, y) + \sum_{(x, y) \in \mathcal{D}_L} \text{KL}(q(e|x, y) || p(e|x, y)) \quad (4)$$

### 3.4 M-step: Explanation-Augmented Prediction model

During M-step, the explanation-augmented prediction model is trained to predict the label of input sentence  $x$  with the explanation  $e$  generated from the variational distribution  $q(e|x, y)$ .

However, note that the label  $y$  is not available during testing, and the explanations for the unlabeled  $x$  can only be generated from the prior distribution  $p(e|x)$ . Therefore, there are some discrepancies between the distributions of the explanations for labeled data in the training stage and those for unlabeled data in the testing stage. To mitigate this issue, in the prediction model, besides sampling an explanation from the variational distribution, we also add a set of explanations from  $p(e|x)$ , which retrieves a set of explanations from similar sentences.

Specifically, given an input sentence  $x$  and the set of labeled and pseudo-labeled data consists of  $(x', e', y')$ , we retrieve  $N$  explanations  $\mathcal{E} := \{e'_i\}_{i=1}^N$  of which the corresponding sentences  $x'$  are the most similar to the input sentence  $x$ , measured by the cosine similarity between the embedding of  $x$  and each  $x'$  from  $\mathcal{D}_E$  under SentenceBERT [19], a pretrained sentence embedding model. Note that we do not directly use a seq2seq model to parametrize  $p(e|x)$  because we find generating explanations without a predicted label often results in irrelevant and even misleading explanations.

---

**Algorithm 1:** Explanation-based Self-Training (ELV-EST)

---

**Input:**  $\mathcal{D}_E = \{(x_1, y_1, e_1), \dots, (x_n, y_n, e_n)\}$ ,  $\mathcal{D}_L = \{(x_{n+1}, y_{n+1}), \dots, (x_m, y_m)\}$ , unlabeled data  $\mathcal{D}_U = \{x_{m+1}, \dots, x_N\}$ , Confidence threshold  $T$

**Output:**  $E_\theta(e|x, y), M_\phi(y|x, E)$

initialize  $E_\theta$  and  $M_\phi$  with  $\mathcal{D}_E \cup \mathcal{D}_L$  using ELV

**repeat**

**for each**  $x_i \in \mathcal{D}_U$  **do**

**if**  $\max_y M_\phi(y|x, E) > T$  **then**

      Assign pseudo-label  $y_i$  to  $x_i$  and generate explanation  $e_i$  with  $E_\theta$

      Update  $\mathcal{D}_L = \mathcal{D}_L \cup (x_i, y_i)$

      Update  $\mathcal{D}_E = \mathcal{D}_E \cup (x_i, y_i, e_i)$  (for explanation retrieval)

      Update  $\mathcal{D}_U = \mathcal{D}_U \setminus x_i$

**end**

**end**

  Train  $E_\theta$  and  $M_\phi$  on  $\mathcal{D}_E \cup \mathcal{D}_L$  with ELV

**until** Convergence or  $\mathcal{D}_U = \emptyset$ 

---

163 Let  $\mathcal{E} = \{e_1, \dots, e_n\}$  denotes all the explanations of  $x$ . For each  $e_i \in \mathcal{E}$ , we feed the explanation  $e_i$   
164 and the input sentence  $x$ , separated by a [SEP] token, to BERT [20] and use the vector at the [CLS]  
165 token to represent the interaction between  $x$  and  $e_i$  as a 768-dimensional feature vector:

$$\mathcal{I}(x, e_i) = \text{BERT}([\text{CLS}]; x; [\text{SEP}]; e_i) \quad (5)$$

166 Our final classifier takes the concatenation of these vectors and outputs the final prediction as:

$$M_\phi(y|x, \mathcal{E}) = \text{MLP}[\text{Average}(\mathcal{I}(x, e_1); \mathcal{I}(x, e_2); \dots; \mathcal{I}(x, e_n))] \quad (6)$$

167 At test time, for each unlabeled  $x$ , we first use  $p(e|x)$  to retrieve a set of explanations and then predict  
168 a label with the explanation-augmented prediction model. Afterward, we can further employ the  
169 explanation generation model to generate an NL explanation to explain the prediction decision based  
170 on both the input sentence and the predicted label.

### 171 3.5 Explanation-based Self-Training

172 As natural language explanations can serve as implicit logic rules, which can generalize to new data  
173 and help assign pseudo-labels to unlabeled data. Therefore, we extend the ELV to the semi-supervised  
174 learning setting and propose an Explanation-based Self-Training (ELV-EST) algorithm. In this  
175 setting, we only have limited labeled examples but abundant unlabeled data  $\mathcal{D}_U = \{x_{m+1}, \dots, x_N\}$ .

176 As illustrated in Algorithm 1, we first use ELV to initialize  $E_\theta$  and  $M_\phi$  with the limited labeled corpus  
177  $\mathcal{D}_E \cup \mathcal{D}_L$ . Afterward, we iteratively use  $M_\phi$  to assign pseudo-labels to unlabeled examples in  $\mathcal{D}_U$  to  
178 extend the labeled data  $\mathcal{D}_L$ . We then use ELV to jointly train  $E_\theta$  and  $M_\phi$  with the augmented labeled  
179 dataset. At the same time, we also employ  $E_\theta$  to generate new explanations with unlabeled examples  
180 and their pseudo-labels. In this way, we can harvest massive pseudo-labels and pseudo-explanations  
181 with unlabeled examples. The pseudo-labeled examples can be used to improve the models while  
182 also enable us to generate more NL explanations. In return, the newly generated explanations can not  
183 only improve the explanation generation model but also serve as implicit rules that help the prediction  
184 model assign more accurate pseudo-labels in the next iteration.

185 The proposed ELV-EST approach is different from the conventional self-training method in two  
186 perspectives. First, in addition to predicting pseudo-labels for unlabeled data, our method also  
187 discovers implicit logic rules in the form of natural language explanations, which in return helps  
188 the prediction model to better assign noisy labels to the unlabeled data. Second, our approach can  
189 produce explainable predictions with  $E_\theta$ . Compared to recent works [7, 13] that parse explanations  
190 to logic forms, our approach does not require task-specific semantic parsers and matching models,  
191 making it task-agnostic and applicable to various natural language understanding tasks with minimal  
192 additional efforts.

Table 1: **Statistics of datasets.** We present the size of train/dev/test sets for 4 datasets in both supervised and semi-supervised settings. Moreover, # Exp means the size of initial explanation sets.

Dataset	# Explanations	# Train (Supervised)	# Train (Semi-supervised)	# Dev	# Test
SemEval [21]	203	7,016	1,210	800	2,715
TACRED [22]	139	68,006	2,751	22,531	15,509
Laptop	70	1,806	135	462	638
Restaurant	75	2,830	107	720	1,120

Table 2: Results (Micro-F1) on Relation Extraction datasets in supervised setting.

Method	TACRED	SemEval
BERT <sub>EM</sub> [25]	66.3	76.9
BERT <sub>EM+MTB</sub> [25]	<b>67.1</b>	77.5
BERT-large	66.4	78.8
BERT-base	64.7	78.3
ELV (M-step only)	65.4	80.2
ELV (ours)	65.9	<b>80.7</b>

Table 3: Results (Macro-F1) on ASC datasets in supervised setting.

Method	Restaurant	Laptop
ASGCN [26]	72.2	71.1
BERT-PT [27]	77.0	75.1
BERT-SPC [28]	77.0	75.0
BERT-base	75.4	72.4
ELV (M-step only)	76.2	74.1
ELV (ours)	<b>77.8</b>	<b>75.2</b>

## 4 Experiments

### 4.1 Datasets

We conduct experiments on two tasks: relation extraction (RE) and aspect-based sentiment classification (ASC). For relation extraction we choose two datasets, TACRED [23] and SemEval [21] in our experiments. We use two customer review datasets, Restaurant and Laptop, which are part of SemEval 2014 Task 4 [24] for the aspect-based sentiment classification task. We use the human-annotated explanations collected in [7] for training our explanation-based models.

### 4.2 Experimental Settings

We conduct experiments in both the **supervised setting** where we have access to all labeled examples in the dataset and the **semi-supervised setting** where we only use a small fraction of labeled examples and considering the rest labeled examples in the original dataset as unlabeled examples by ignoring their labels. In both settings, only a few human-annotated NL explanations are available. The number of explanations, labeled data used in supervised/unsupervised setting, and the statistics of the datasets are presented in Table 1.

We employ BERT-base and UniLM-base as the backbone of our prediction model and explanation generation model, respectively. We select batch size over  $\{32, 64\}$  and learning rate over  $\{1e-5, 2e-5, 3e-5\}$ . The number of retrieved explanations is set to 10 for all tasks. We train the prediction model for 3 epochs and the generation model for 5 epochs in each EM iteration. We use Adam optimizers and early stopping with the best validation F1-score.

### 4.3 Compared Methods

In the **supervised setting**, we compare ELV with the BERT-base baseline that directly fine-tunes the pre-trained BERT-base model on the target datasets. To show the importance of modeling the interaction between the explanation generation model and the explanation-augmented prediction model, we also compare with a variant of our model, which only trains the explanation-augmented prediction module with all the explanations generated from the prior distribution, denoted as ELV (M-step only)). We also compare with some state-of-the-art algorithms on the RE and SA tasks.

In the **semi-supervised setting**, we compare ELV-EST against several competitive semi-supervised text classification methods including Pseudo-Labeling [29], Self-Training [30], and Data Programming [13] which incorporates NL explanations to perform semi-supervised text classification. Note that all compared model variants incorporate BERT-base as the backbone model.



Table 4: Results (Micro-F1) on Relation Extraction datasets in semi-supervised setting.

Method	TACRED	SemEval
BERT-base	25.1	49.3
Pseudo-Labeling [29]	28.6	50.2
Self-Training [30]	36.9	59.5
Data Programming [13]	25.8	47.9
ELV-EST (ours)	<b>42.5</b>	<b>66.4</b>

Table 5: Results (Macro-F1) on ASC datasets in semi-supervised setting.

Method	Restaurant	Laptop
BERT-base	32.2	34.6
Pseudo-Labeling [29]	42.5	38.2
Self-Training [30]	47.2	42.3
Data Programming [13]	38.2	36.3
ELV-EST (ours)	<b>59.5</b>	<b>63.6</b>

Table 6: Human evaluation results. The scores scale from 1 to 5 (the larger, the better). The inner-rater agreement measured by Kappa score is 0.51.

Model	Inf.	Corr.	Cons.
Seq2Seq	2.43	3.27	2.68
Transformer	2.35	3.12	2.62
UniLM	3.48	3.94	3.14
ELV (ours)	<b>3.87</b>	<b>4.20</b>	<b>3.51</b>

Table 7: Results on ASC datasets with explanations with words randomly corrupted (80%). Orig + Rand Exp is the 1:1 mix of original and randomly corrupted explanations.

Method	Restaurant	Laptop
BERT-base	75.4	72.4
w. 80% Rand Word	73.2	70.9
Orig + Rand Exp	76.9	74.0
ELV (ours)	<b>77.8</b>	<b>75.2</b>

## 223 4.4 Experimental Results

224 **Results on supervised setting.** We first present the results in the supervised setting in Table 2  
 225 and 3. ELV significantly outperforms the strong BERT baseline in all four datasets, demonstrating  
 226 the effectiveness of exploiting NL explanations as additional information for natural language  
 227 understanding. ELV also consistently outperforms the ELV (M-step only), showing that ELV’s  
 228 variational EM training effectively models the interactions between explanation and prediction. Also,  
 229 the performance of ELV compares favorably against several competitive recent studies focusing on  
 230 RE and ASC respectively, further demonstrating the effectiveness of ELV.

231 **Results on semi-supervised setting.** The results in the semi-supervised setting are presented in  
 232 Table 4 and 5. In the semi-supervised scenario, ELV-EST method significantly outperforms various  
 233 semi-supervised text classification methods, as well as the data programming approach. The latter  
 234 uses pre-defined rules to parse the NL explanations into logic forms and match unlabeled examples,  
 235 on all four datasets. The improvement upon the BERT-base + self-training baseline is around 7 points  
 236 for RE datasets and over 15 points for ASC datasets in terms of F1 score. This demonstrates the  
 237 effectiveness of ELV-EST in the semi-supervised setting.

238 **Results on explanation generation.** We further evaluate the quality of the explanation generation  
 239 model with human evaluation. We invite 5 graduate students with English proficiency to score the  
 240 explanations generated on the test set with input sentence and the label predicted by the explanation-  
 241 augmented prediction module<sup>1</sup>. The annotation scenarios include the explanations’ informativeness  
 242 (Info.), correctness (Corr.), and consistency (Cons.) with respect to the model prediction. The  
 243 inner-rater agreement is at 0.51 Kappa score. The details of human evaluation and examples of  
 244 generated explanations are presented in the Appendix due to space constraints.

245 For comparison, we include a fine-tuned a UniLM model with annotated NL explanations, as well as  
 246 two baselines trained from scratch using annotated NL explanations, one with a vanilla transformer  
 247 model and the other an attention-based LSTM seq2seq model. The results are in Table 6. The  
 248 explanations generated by our ELV framework are substantially better than those generated by the  
 249 fine-tuned UniLM model. ELV generates better NL explanations that are relevant to the model’s  
 250 decision-making process, because it models the interaction of the explanation generation model and  
 251 the prediction model.

<sup>1</sup>The prediction module is jointly trained with the explanation generation module




Figure 2: Performance with different number of explanations. We compare our method with baseline(s) in both supervised setting (left) and semi-supervised setting (right).




Figure 3: Performance with different number of labeled or unlabeled data in supervised setting (left) and semi-supervised setting (right) respectively.

## 4.5 Analysis

**Performance with corrupted explanations.** We first investigate the model performance w.r.t. the quality of the retrieved explanations. We compare with corrupted explanations which randomly replace 80% percentages of the words in the original explanations, results shown in Table 7. The performance with corrupted explanations significantly decreases as expected. The high-quality explanations help the model better generalize while the random ones may confuse the model.

**Performance with different numbers of explanations.** We then investigate the performances with different amounts of explanations. As illustrated in Figure 2 (left), with as few as 15 annotated explanations, ELV significantly outperforms its counterpart trained without the variational EM framework in the supervised setting. The performance of ELV continues to improve with more explanations but the performance of ELV (M-step only) starts to saturate with 45 explanations, showing the importance of modeling the interaction between the explanation generation model and explanation-augmented prediction model. Similar results are observed in the semi-supervised learning setting in Figure 2 (right).

**Performance with different numbers of labeled/unlabeled data.** We also investigate the performance of different models with different proportions of training data. From Figure 3 (left), we can see that ELV consistently outperforms the BERT baseline with different amounts of labeled data. Especially, the improvement is the most significant when only 10% of labeled data is used. This is because human explanations provide additional supervision and can serve as implicit logic rules to help generalization. In the semi-supervised learning setting (Figure 3, right), ELV-EST outperforms traditional self-training methods by a large margin, especially with fewer unlabeled data, further confirming the improved generalization ability from explanations.

## 5 Conclusion

In this paper, we propose ELV, a novel framework for training interpretable natural language understanding models with limited human annotated explanations. Our framework treats natural language explanations as latent variables that model the underlying reasoning process to enable interactions between explanation generation and explanation-based classification. Experimental results in both supervised and semi-supervised settings show that ELV is not only able to make effective predictions but also generate meaningful explanations. In the future, we plan to apply our framework to other natural language understanding tasks.



## Broader Impact

Deep learning has achieved great success in natural language understanding. However, most existing systems are not interpretable, which limit their applications to many domains such as healthcare, finance, and legislation. In these domains, interpretability is a high priority. This paper proposed a principled probabilistic model for text classification, which not only makes effective prediction but also offers good explainability. Though the model is developed for the task of text classification, it is a very general framework and could be generalized to other tasks in natural language understanding. Such a system could be useful in a variety of tasks such as decision making with clinical notes in healthcare, justice, and criminal identification with legal data, and risk management in finance.

On the other hand, such a system also brings potential risks depending on the quality of the generated natural language explanations. For example, the generated natural language could have certain biases, which have been reported in many natural language understanding systems [31, 32]. How to mitigate these risks will be our future work.

## References

- [1] Braden Hancock, Paroma Varma, Stephanie Wang, Martin Bringmann, Percy Liang, and Christopher Ré. Training classifiers with natural language explanations. In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1884–1895, Melbourne, Australia, July 2018. Association for Computational Linguistics. doi: 10.18653/v1/P18-1175. URL <https://www.aclweb.org/anthology/P18-1175>.
- [2] Nazneen Fatema Rajani, Bryan McCann, Caiming Xiong, and Richard Socher. Explain yourself! leveraging language models for commonsense reasoning. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 4932–4942, Florence, Italy, July 2019. Association for Computational Linguistics. doi: 10.18653/v1/P19-1487. URL <https://www.aclweb.org/anthology/P19-1487>.
- [3] Oana-Maria Camburu, Tim Rocktäschel, Thomas Lukasiewicz, and Phil Blunsom. e-snli: Natural language inference with natural language explanations. In *Advances in Neural Information Processing Systems*, pages 9539–9549, 2018.
- [4] Nazneen Fatema Rajani, Bryan McCann, Caiming Xiong, and Richard Socher. Explain yourself! leveraging language models for commonsense reasoning. *arXiv preprint arXiv:1906.02361*, 2019.
- [5] Shashank Srivastava, Igor Labutov, and Tom Mitchell. Joint concept learning and semantic parsing from natural language explanations. In *Proceedings of the 2017 Conference on Empirical Methods in Natural Language Processing*, pages 1527–1536, Copenhagen, Denmark, September 2017. Association for Computational Linguistics. doi: 10.18653/v1/D17-1161. URL <https://www.aclweb.org/anthology/D17-1161>.
- [6] Wenxuan Zhou, Hongtao Lin, Bill Yuchen Lin, Ziqi Wang, Junyi Du, Leonardo Neves, and Xiang Ren. Nero: A neural rule grounding framework for label-efficient relation extraction. In *Proceedings of The Web Conference 2020, WWW ’20*, page 2166–2176, New York, NY, USA, 2020. Association for Computing Machinery. ISBN 9781450370233. doi: 10.1145/3366423.3380282. URL <https://doi.org/10.1145/3366423.3380282>.
- [7] Ziqi Wang, Yujia Qin, Wenxuan Zhou, Jun Yan, Qinyuan Ye, Leonardo Neves, Zhiyuan Liu, and Xiang Ren. Learning from explanations with neural execution tree. In *International Conference on Learning Representations*, 2020. URL <https://openreview.net/forum?id=rJlUtOEYwS>.
- [8] Jason Palmer, Kenneth Kreutz-Delgado, Bhaskar D Rao, and David P Wipf. Variational em algorithms for non-gaussian latent variable models. In *Advances in neural information processing systems*, pages 1059–1066, 2006.
- [9] Shashank Srivastava, Igor Labutov, and Tom Mitchell. Joint concept learning and semantic parsing from natural language explanations. In *Proceedings of the 2017 conference on empirical methods in natural language processing*, pages 1527–1536, 2017.

- [10] Shikhar Murty, Pang Wei Koh, and Percy Liang. Expbert: Representation engineering with natural language explanations. *arXiv preprint arXiv:2005.01932*, 2020.
- [11] Dan Goldwasser and Dan Roth. Learning from natural instructions. *Machine learning*, 94(2): 205–232, 2014.
- [12] Sanja Fidler et al. Teaching machines to describe images with natural language feedback. In *Advances in Neural Information Processing Systems*, pages 5068–5078, 2017.
- [13] Braden Hancock, Martin Bringmann, Paroma Varma, Percy Liang, Stephanie Wang, and Christopher Ré. Training classifiers with natural language explanations. In *Proceedings of the conference. Association for Computational Linguistics. Meeting*, volume 2018, page 1884. NIH Public Access, 2018.
- [14] Diederik P Kingma and Max Welling. Auto-encoding variational bayes. *arXiv preprint arXiv:1312.6114*, 2013.
- [15] Zied Bouraoui, Jose Camacho-Collados, and Steven Schockaert. Inducing relational knowledge from bert. *arXiv preprint arXiv:1911.12753*, 2019.
- [16] Fabio Petroni, Tim Rocktäschel, Patrick Lewis, Anton Bakhtin, Yuxiang Wu, Alexander H Miller, and Sebastian Riedel. Language models as knowledge bases? *arXiv preprint arXiv:1909.01066*, 2019.
- [17] Adam Roberts, Colin Raffel, and Noam Shazeer. How much knowledge can you pack into the parameters of a language model? *arXiv preprint arXiv:2002.08910*, 2020.
- [18] Li Dong, Nan Yang, Wenhui Wang, Furu Wei, Xiaodong Liu, Yu Wang, Jianfeng Gao, Ming Zhou, and Hsiao-Wuen Hon. Unified language model pre-training for natural language understanding and generation. In *Advances in Neural Information Processing Systems*, pages 13042–13054, 2019.
- [19] Nils Reimers and Iryna Gurevych. Sentence-BERT: Sentence embeddings using Siamese BERT-networks. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 3982–3992, Hong Kong, China, November 2019. Association for Computational Linguistics. doi: 10.18653/v1/D19-1410. URL <https://www.aclweb.org/anthology/D19-1410>.
- [20] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. BERT: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 4171–4186, Minneapolis, Minnesota, June 2019. Association for Computational Linguistics. doi: 10.18653/v1/N19-1423. URL <https://www.aclweb.org/anthology/N19-1423>.
- [21] Iris Hendrickx, Su Nam Kim, Zornitsa Kozareva, Preslav Nakov, Diarmuid O Séaghdha, Sebastian Padó, Marco Pennacchiotti, Lorenza Romano, and Stan Szpakowicz. Semeval-2010 task 8: Multi-way classification of semantic relations between pairs of nominals. In *Proceedings of the 5th International Workshop on Semantic Evaluation*, pages 33–38. Association for Computational Linguistics, 2010.
- [22] Yuhao Zhang, Victor Zhong, Danqi Chen, Gabor Angeli, and Christopher D. Manning. Position-aware attention and supervised data improve slot filling. In *Proceedings of the 2017 Conference on Empirical Methods in Natural Language Processing (EMNLP 2017)*, pages 35–45, 2017. URL <https://nlp.stanford.edu/pubs/zhang2017tacred.pdf>.
- [23] Yuhao Zhang, Victor Zhong, Danqi Chen, Gabor Angeli, and Christopher D Manning. Position-aware attention and supervised data improve slot filling. In *Proceedings of the 2017 Conference on Empirical Methods in Natural Language Processing*, pages 35–45, 2017.

- 379 [24] Maria Pontiki, Dimitris Galanis, John Pavlopoulos, Harris Papageorgiou, Ion Androutsopoulos,  
380 and Suresh Manandhar. SemEval-2014 task 4: Aspect based sentiment analysis. In *Proceedings*  
381 *of the 8th International Workshop on Semantic Evaluation (SemEval 2014)*, pages 27–35, Dublin,  
382 Ireland, August 2014. Association for Computational Linguistics. doi: 10.3115/v1/S14-2004.  
383 URL <https://www.aclweb.org/anthology/S14-2004>.
- 384 [25] Livio Baldini Soares, Nicholas FitzGerald, Jeffrey Ling, and Tom Kwiatkowski. Matching  
385 the blanks: Distributional similarity for relation learning. In *Proceedings of the 57th Annual*  
386 *Meeting of the Association for Computational Linguistics*, pages 2895–2905, Florence, Italy,  
387 July 2019. Association for Computational Linguistics. doi: 10.18653/v1/P19-1279. URL  
388 <https://www.aclweb.org/anthology/P19-1279>.
- 389 [26] Chen Zhang, Qiuchi Li, and Dawei Song. Aspect-based sentiment classification with aspect-  
390 specific graph convolutional networks. *arXiv preprint arXiv:1909.03477*, 2019.
- 391 [27] Hu Xu, Bing Liu, Lei Shu, and Philip S Yu. Bert post-training for review reading comprehension  
392 and aspect-based sentiment analysis. *arXiv preprint arXiv:1904.02232*, 2019.
- 393 [28] Youwei Song, Jiahai Wang, Tao Jiang, Zhiyue Liu, and Yanghui Rao. Attentional encoder  
394 network for targeted sentiment classification. *arXiv preprint arXiv:1902.09314*, 2019.
- 395 [29] Dong-Hyun Lee. Pseudo-label: The simple and efficient semi-supervised learning method for  
396 deep neural networks. In *Workshop on challenges in representation learning, ICML*, volume 3,  
397 page 2, 2013.
- 398 [30] Chuck Rosenberg, Martial Hebert, and Henry Schneiderman. Semi-supervised self-training of  
399 object detection models. *WACV/MOTION*, 2, 2005.
- 400 [31] Jieyu Zhao, Tianlu Wang, Mark Yatskar, Vicente Ordonez, and Kai-Wei Chang. Men also like  
401 shopping: Reducing gender bias amplification using corpus-level constraints. *arXiv preprint*  
402 *arXiv:1707.09457*, 2017.
- 403 [32] Marta R Costa-jussà. An analysis of gender bias studies in natural language processing. *Nature*  
404 *Machine Intelligence*, pages 1–2, 2019.

## 405 A Examples of Explanations

406 In this section we present several human annotated natural language explanations and explanationos  
407 generated by the explanation generation model in the ELV framework in different datasets to better  
408 understand the proposed approach.

### 409 A.1 TACRED

#### 410 A.1.1 Human Annotated Explanations

411 a) *Although not a Playboy Playmate , she has appeared in nude pictorials with her Girls Next*  
412 *Door costars and fellow Hefner girlfriends Holly Madison and Kendra Baskett(OBJ), then*  
413 *known as Kendra Wilkinson(SUBJ).*

414 **Label:** per:alternate\_names

415 **Explanation:** The term "then known as" occurs between SUBJ and OBJ and there are no  
416 more than six words between SUBJ and OBJ

417 b) *Burke(SUBJ) 's mother Melissa Bell(OBJ) was a singer in the dance group Soul II Soul,*  
418 *which had hits in the 1980s and 1990s.*

419 **Label:** per:parents

420 **Explanation:** SUBJ and OBJ sandwich the phrase "'s mother" and there are no more than  
421 three words between SUBJ and OBJ

422 c) *Ellen Pompeo(OBJ) secretly married Chris Ivery(SUBJ) Congratulations to the newlyweds*  
423 *and let them live happily ever after !!*

424 **Label:** per:spouse

425 **Explanation:** There are no more than four words between SUBJ and OBJ and SUBJ and  
426 OBJ sandwich the phrase "secretly married"

#### 427 A.1.2 Machine Generated Explanations

428 a) *Lomax shares a story about Almena Lomax, his mother and a newspaper owner and*  
429 *journalist(OBJ) in Los Angeles , taking her(SUBJ) family on the bus to Tuskegee, Ala., in*  
430 *1961.*

431 **Label:** per:title

432 **Explanation:** the word "family" is right after SUBJ.

433 b) *What happened to their investments was of no interest to them, because they would already be*  
434 *paid , said Paul Hodgson(OBJ), senior research associate at the Corporate Library(SUBJ),*  
435 *a shareholder activist group.*

436 **Label:** per:org:top\_members/employees

437 **Explanation:** The word "senior research associate at" appears right before SUBJ.

438 c) *Iroquois passport dispute raises sovereignty issue The National Congress of American*  
439 *Indians(SUBJ), based in Washington, DC(OBJ), has advocated on behalf of the lacrosse*  
440 *team, urging British officials to allow the members entry into England on their Iroquois-*  
441 *issued passports.*

442 **Label:** org:country\_of\_headquarterse

443 **Explanation:** The word "based in" appears right before OBJ.

### 444 A.2 SemEval

#### 445 A.2.1 Human Annotated Explanations

446 a) *Morton's SUBJ-O is the most common cause of localized OBJ-O in the third interspace and*  
447 *these diagnostic tests produce good indications of the condition.*

448 **Label:** Cause-Effect(e1,e2)

449 **Explanation:** Between SUBJ and OBJ the term "is the most common cause of" appears  
450 and SUBJ precedes OBJ

- 451 b) *The frontal SUBJ-O is a part of the OBJ-O that maintains very close ties with the limbic*  
 452 *system.*  
 453 **Label:** Component-Whole(e1,e2)  
 454 **Explanation:** Between SUBJ and OBJ the term "is a part of the" appears and SUBJ precedes  
 455 OBJ.  
 456 c) *Out current Secretary is gathering SUBJ-O from past OBJ-O and committee chairs.*  
 457 **Label:** Entity-Origin(e1,e2)  
 458 **Explanation:** The phrase "from past" links SUBJ and OBJ and there are no more than three  
 459 words between SUBJ and OBJ and OBJ follows SUBJ.

## 460 A.2.2 Machine Generated Explanations

- 461 a) *SUBJ-O caused OBJ-O at the Charlotte Douglas International Airport Monday morning.*  
 462 **Label:** Cause-Effect(e1,e2)  
 463 **Explanation:** There is only one word "caused" between subj and obj and obj follows subj.  
 464 b) *The base in which it sits hides the damage which occurred when the SUBJ-O was removed*  
 465 *from its initial OBJ-O on the Long Tan battlefield.*  
 466 **Label:** Entity-Origin(e1,e2)  
 467 **Explanation:** Between subj and obj the phrase "was removed into" occurs and there are no  
 468 more than four words between SUBJ and OBJ and SUBJ precedes OBJ  
 469 c) *A hinge SUBJ-O attaches a OBJ-O pivotally to a base of an electronic device and has a*  
 470 *pivoting leaf and a stationary leaf.*  
 471 **Label:** Component-Whole(e1,e2)  
 472 **Explanation:** The phrase "attaches a" between SUBJ and OBJ and OBJ follows SUBJ

## 473 A.3 Restaurant

### 474 A.3.1 Human Annotated Explanations

- 475 a) *We had great desserts (including the best cannoli I've ever had) and then they offered an*  
 476 *after dinner drink, on the house. (Term: cannoli)*  
 477 **Label:** positive  
 478 **Explanation:** The word "best" directly precedes the term.  
 479 b) *All the desserts the group tried got favorable reviews. (Term: desserts)*  
 480 **Label:** positive  
 481 **Explanation:** The string "favorable" appears no more than 5 words after the term.  
 482 c) *The most annoying thing, though, is the fact that the servers seem to be trained to drive*  
 483 *revenue. (Term: servers)*  
 484 **Label:** negative  
 485 **Explanation:** The word "annoying" occurs before the term.

### 486 A.3.2 Machine Generated Explanations

- 487 a) *This little place is wonderfully warm welcoming. (Term: place)*  
 488 **Label:** positive  
 489 **Explanation:** The term is followed by "wonderful".  
 490 b) *The falafal was rather over cooked and dried but the chicken was fine. (Term: chicken)*  
 491 **Label:** positive  
 492 **Explanation:** The word "fine" occurs within 3 words after the term.  
 493 c) *Service was awful - mostly because staff were overwhelmed on a Saturday night. (Term:*  
 494 *staff)*  
 495 **Label:** negative  
 496 **Explanation:** The word "unbearable" occurs within three words after the term.

Table 8: Results (Macro-F1) on ASC datasets in supervised setting with different fraction of labeled data used.

Fraction of labeled data used	60%	70%	80%	90%	100%
ELV (ours)	75.5	-	-	-	-
BERT-base	74.6	75.0	75.3	75.4	75.4

## 497 A.4 Laptop

### 498 A.4.1 Human Annotated Explanations

- 499 a) *The DVD drive randomly pops open when it is in my backpack as well, which is annoying.*  
500 (Term: DVD drive)  
501 **Label:** negative  
502 **Explanation:** The string "annoying" occurs after the term
- 503 b) *The Apple team also assists you very nicely when choosing which computer is right for you.*  
504 (Term: Apple team)  
505 **Label:** positive  
506 **Explanation:** The string "very nicely" occurs after the term by no more than 6 words.
- 507 c) *The design is awesome, quality is unprecedented. (Term: design)*  
508 **Label:** positive  
509 **Explanation:** The word "awesome" is within 2 words after the term.

### 510 A.4.2 Machine Generated Explanations

- 511 a) *I ordered my 2012 mac mini after being disappointed with spec of the new 27 Imacs. (Term:*  
512 *spec)*  
513 **Label:** negative  
514 **Explanation:** The word "disappointed" occurs within 3 words before the term.
- 515 b) *I found the mini to be exceptionally easy to set up. (Term: set up)*  
516 **Label:** positive  
517 **Explanation:** The phrase "exceptionally easy" occurs within 3 words before the term.
- 518 c) *However, there are MAJOR issues with the touchpad which render the device nearly useless.*  
519 *(Term: touchpad)*  
520 **Label:** negative  
521 **Explanation:** The phrase "nearly useless" occurs within 3 words after the term.

## 522 B Human Evaluation Details

523 For human evaluation, we random sample 100 examples in the test set of the Restaurant dataset and  
524 use ELV to predict the labels of the selected examples. Then we use different compared models  
525 to generate explanations of the prediction results with both the input sentences and the predicted  
526 labels. Afterward, we invite 5 graduate students with English proficiency to score the explanations.  
527 The annotation scenarios include the explanantions' informativeness(Info.), correctness (Corr.), and  
528 consistency (Cons.) with respect to the model prediction. Specifically, the informativeness measures  
529 to what extent the generated explanation is helpful to understand the model's prediction output. The  
530 correctness measures whether the explanation is factually correct (e.g. word "good" before the terms  
531 leads to positive label while word "annoying" is negative). The consistency refers to whether the  
532 explanation is consistent with the input sentence (i.e. the description in the explanation is true w.r.t  
533 the input sentence).



## 534 C Addition Analysis Experiments

535 In this section, we report additional experimental results comparing the performance of ELV with  
536 60% of labeled data and the performance with the BERT-base model with 60%, 70%, 80%, 90%, and  
537 100% of labeled data to investigate to what extent human annotated explanations can replace human  
538 labeled examples. The result is shown in Table 8. We find that ELV can achieve or even exceed the  
539 performance of a BERT-base model trained with much more labeled data. This confirms that ELV  
540 can effectively leverage human annotated explanations as additional information.